

Map Reduce Design Patterns

Caveat

- The great power of MapReduce lays in its simplicity
 - The programmer only needs to prepare input data, configuration information, implement mapper and reducer and, optionally, setup, cleanup, combiner and partitioner
 - All other aspects are handled by the framework
- This simplicity has some pros and cons
 - Pros -- It provides clear boundaries for what you can do and cannot do
 - Cons -- The programmer has to express any algorithm in terms of a small number of rigidly defined components that must fit together in very specific ways
- How to recast/reshape an algorithm into this programming model, i.e., as a sequence of one or more MapReduce jobs?
- Good programmers can produce bad MapReduce algorithms
- MapReduce design patterns
 - Principles/guidelines to build better software
 - Not specific to a domain (e.g., text processing). Instead, it is a general approach to solve a problem

Design Patterns

- In-mapper Combining
- Pairs and stripes
- Patterns for relational algebra operators
- Patterns for matrix-vector and matrix-matrix multiplication

In-mapper Combining (I)

Intermediate data (i.e., mappers outputs) are

- Written locally
- Transferred from mappers to reducers over cluster network
- Issue
 - Bottleneck: Writing to disk and transferring over the network are relatively «expensive» operations
 - Reducing intermediate data translates into increasing algorithmic efficiency
- Solution: Local aggregation
 - Use combiners
 - Use In-Mapper Combining

In-Mapper Combining (II)

Pseudo-code of the basic word-count algorithm

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $t \in \text{doc } d$  do
4:       EMIT(term  $t$ , count 1)

1: class REDUCER
2:   method REDUCE(term  $t$ , counts  $[c_1, c_2, \dots]$ )
3:      $sum \leftarrow 0$ 
4:     for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
5:        $sum \leftarrow sum + c$ 
6:     EMIT(term  $t$ , count  $sum$ )
```

In-Mapper Combining (III)

Pseudo-code of stateless in-mapper combining

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:      $H \leftarrow$  new ASSOCIATIVEARRAY
4:     for all term  $t \in$  doc  $d$  do
5:        $H\{t\} \leftarrow H\{t\} + 1$ 
6:     for all term  $t \in H$  do
7:       EMIT(term  $t$ , count  $H\{t\}$ )
```

In-Mapper Combining (IV)

Pseudo-code of stateful in-mapper combining (the true in-mapper combining)

```
1: class MAPPER
2:   method INITIALIZE
3:      $H \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:   method MAP(docid  $a$ , doc  $d$ )
5:     for all term  $t \in \text{doc } d$  do
6:        $H\{t\} \leftarrow H\{t\} + 1$ 
7:   method CLOSE
8:     for all term  $t \in H$  do
9:       EMIT(term  $t$ , count  $H\{t\}$ )
```


In-Mapper Combining (V)

- Advantages:

- Combiners are provided as an optimization to the execution framework, which has the option of using it (perhaps multiple times) or not at all. On the other hand, in-mapper combining provides complete control on local aggregation process (how and when)
- Combiners only reduce the amount of intermediate data that is transferred across the network but do not reduce intermediate data emitted by mappers in the first place. In-mapper combining provides further control by avoiding unnecessary object creation as output to the mapper

- Disadvantages:

- Potential ordering-dependent bugs – preserving state across records may cause algorithmic behavior to depend on the order in which input records are encountered
- Memory scalability bottleneck -- the associative array may no longer fit in memory. Solved by:
 - Either emitting intermediate data after processing every n (how big?) key-value pairs rather than all of them
 - Or monitoring memory footprint and flush intermediate data to disk once memory usage has exceeded a certain (how big?) threshold

Matrix Generation

- Common problem:
 - Given an input of size N , generate a squared output of size $N \times N$
- Example: word co-occurrence matrix
 - Given a document collection, emit the bigram frequencies
 - Where N is the number of unique words
 - Cell m_{ij} contains the number of times word w_i co-occurs with word w_j
 - Used for example in statistical natural language processing
- Two solutions
 - Pairs: generating $O(N^2)$ data in $O(1)$ space
 - Stripes: generation $O(N)$ data in $O(N)$ space

Pairs

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $w \in \text{doc } d$  do
4:       for all term  $u \in \text{NEIGHBORS}(w)$  do
5:         EMIT(pair ( $w, u$ ), count 1)                                ▷ Emit count for each co-occurrence

1: class REDUCER
2:   method REDUCE(pair  $p$ , counts [ $c_1, c_2, \dots$ ])
3:      $s \leftarrow 0$ 
4:     for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
5:        $s \leftarrow s + c$                                             ▷ Sum co-occurrence counts
6:     EMIT(pair  $p$ , count  $s$ )
```

Stripes

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $w \in \text{doc } d$  do
4:        $H \leftarrow \text{new ASSOCIATIVEARRAY}$ 
5:       for all term  $u \in \text{NEIGHBORS}(w)$  do
6:          $H\{u\} \leftarrow H\{u\} + 1$  ▷ Tally words co-occurring with  $w$ 
7:       EMIT(Term  $w$ , Stripe  $H$ )

1: class REDUCER
2:   method REDUCE(term  $w$ , stripes [ $H_1, H_2, H_3, \dots$ ])
3:      $H_f \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:     for all stripe  $H \in \text{stripes } [H_1, H_2, H_3, \dots]$  do
5:       SUM( $H_f, H$ ) ▷ Element-wise sum
6:     EMIT(term  $w$ , stripe  $H_f$ )
```

Relational Algebra Operators

Use case:
Search engines

A_1	A_2	A_3	A_4	A_5	<i>Schema</i>
					<i>Tuple t</i>

Relation R

- **SELECTION:** Select from relation R tuples satisfying condition $c(t)$
- **PROJECTION:** For each tuple in relation R , select only certain attributes A_i
- **UNION, INTERSECTION, DIFFERENCE:** Set operations on two relations with same schema
- **NATURAL JOIN**
- **GROUPING and AGGREGATION**

Selection and projection

Selection

- Map: each tuple t in R , if condition $c(t)$ is satisfied, is outputted as a $(t, \perp)$ pair
- Reduce: for each $(t, \perp, \perp, \perp, \dots)$ pair in input, output $(t, \perp)$

Projection

- Map: each tuple t in R , create a new tuple t' containing only the projected attributes and output a $(t', \perp)$ pair
- Reduce: for each $(t', \perp, \perp, \perp, \dots)$ pair in input, output $(t', \perp)$

Union, intersection and difference

Union

- Map: for each tuple t in R and S , output a $(t, \perp)$ pair
- Reduce: for each input key t , there will be 1 or 2 values equal to $\perp$. Coalesce them in a single output $(t, \perp)$

Intersection

- Map: for each tuple t in R and S , output a $(t, \perp)$ pair
- Reduce: for each input key t , there will be 1 or 2 values equal to $\perp$. If there are 2 value, coalesce them in a single output $(t, \perp)$, otherwise do nothing

Difference

- Map: for each tuple t in R , output (t, R) and for each tuple t in S , output (t, S)
- Reduce: for each input key t , there will be 1 or 2 values. If there is 1 value equal to (t, R) output $(t, \perp)$, otherwise do nothing

Natural Join

For simplicity, assume we have two relations $R(A,B)$ and $S(B,C)$. Find tuples that agree on the B attribute values and output them.

Natural join

- Map: for each tuple (a, b) from R , output $(b, (\text{R}, a))$ and for each tuple (b, c) from S , produce $(b, (\text{S}, c))$
- Reduce: For each input key b , there will be a list of values of the form (R, a) or (S, c) . Construct all pairs and output them together with b

Grouping and aggregation

For simplicity, assume we have the relation $R(A,B,C)$ and we want to group-by A and aggregate on B , disregarding C .


Grouping and aggregation

- Map: for each tuple (a, b, c) from R , output (a, b) . Each key a represents a group.
- Reduce: apply the aggregation operator to the list of b values associated with group keyed by a , producing x . Then output (a, x) .

Stage Chaining

- As map reduce calculations get more complex, it's useful to break them down into stages, with the output of one stage serving as input to the next

Matrix-Vector Multiplication



The matrix does not fit in memory, and

1. The vector v does fit in a machine's memory
2. The vector v does not fit in machine's memory

A graphical and simple example of Matrix Vector Multiplication


$$\begin{bmatrix} 1 & 5 \\ 2 & 4 \\ 3 & 6 \end{bmatrix} \times \begin{bmatrix} 7 \\ 2 \end{bmatrix} = \begin{bmatrix} 17 \\ 22 \\ 33 \end{bmatrix}$$

Matrix M
size 3×2

Vector v
size 2×1

Vector r
size 3×1


Vector does fit



The matrix is stored in HDFS as a list of (i, j, m_{ij}) tuples

- The elements v_j of v are available to all mappers
- Map: $((i, j), m_{ij})$ pair $\rightarrow (i, m_{ij}v_j)$ pair
- Reduce: $(i, [m_{i1}v_1, m_{i2}v_2, \dots, m_{in}v_n])$ pair $\rightarrow (i, m_{i1}v_1 + m_{i2}v_2 + \dots + m_{in}v_n)$ pair

Vector does not fit




Matrix M
size $m \times n$



$\times$

Vector v
size $n \times 1$

Vector does not fit




Matrix M
size $m \times n$



$\times$

Vector v
size $n \times 1$

Vector does not fit




Matrix M
size $m \times n$



$\times$


Vector v
size $n \times 1$

Vector does not fit



- Divide the vector in equal-sized subvectors that can fit in memory
- According to that, divide the matrix in stripes
- Stripe i and subvector i are independent from other stripes/subvectors
- Use the previous algorithm for each stripe/subvector pair

Matrix-Matrix Multiplication (I)




Matrix M

Matrix N

Matrix P

$$p_{ik} = \sum_j m_{ij} n_{jk}$$

Matrix-Matrix Multiplication (I)



Matrix M

Matrix N

Matrix P

$$p_{ik} = \sum_j m_{ij} n_{jk}$$

Matrix-Matrix Multiplication (II)

- A matrix can be seen as a 4-attributes relation/table:
 - (Matrix name, row index, column index, value) tuples
 - $M \rightarrow (M, i, j, m_{ij}), N \rightarrow (N, j, k, n_{jk})$
- As large matrices are often sparse (0's) we omit such tuples

Matrix-Matrix Multiplication (III)

Algorithm 1: The Map Function

```
1 for each element  $m_{ij}$  of  $M$  do
2   produce (key, value) pairs as  $((i, k), (M, j, m_{ij}))$ , for  $k = 1, 2, 3, \dots$  up
   to the number of columns of  $N$ 
3 for each element  $n_{jk}$  of  $N$  do
4   produce (key, value) pairs as  $((i, k), (N, j, n_{jk}))$ , for  $i = 1, 2, 3, \dots$  up
   to the number of rows of  $M$ 
5 return Set of (key, value) pairs that each key,  $(i, k)$ , has a list with
   values  $(M, j, m_{ij})$  and  $(N, j, n_{jk})$  for all possible values of  $j$ 
```

Algorithm 2: The Reduce Function

```
1 for each key  $(i, k)$  do
2   sort values begin with  $M$  by  $j$  in  $list_M$ 
3   sort values begin with  $N$  by  $j$  in  $list_N$ 
4   multiply  $m_{ij}$  and  $n_{jk}$  for  $j_{th}$  value of each list
5   sum up  $m_{ij} * n_{jk}$ 
6 return  $(i, k), \sum_{j=1} m_{ij} * n_{jk}$ 
```

Matrix Matrix Multiplication (IV)

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} a & b \\ c & d \\ e & f \end{bmatrix} \longrightarrow \begin{bmatrix} 1a + 2c + 3e & 1b + 2d + 3f \\ 4a + 5c + 6e & 4b + 5d + 6f \end{bmatrix}$$

Matrix Matrix Multiplication (IV)

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} a & b \\ c & d \\ e & f \end{bmatrix} \longrightarrow \begin{bmatrix} 1a + 2c + 3e & 1b + 2d + 3f \\ 4a + 5c + 6e & 4b + 5d + 6f \end{bmatrix}$$

$(i, k), (M, j, m_{ij})$

$$m_{11} = 1$$

$$(1, 1), (M, 1, 1)k = 1$$

$$(1, 2), (M, 1, 1)k = 2$$

$$m_{12} = 2$$

$$(1, 1), (M, 2, 2)k = 1$$

$$(1, 2), (M, 2, 2)k = 2$$

.....

$$m_{23} = 6$$

$$(2, 1), (M, 3, 6)k = 1$$

$$(2, 2), (M, 3, 6)k = 2$$

$(i, k), (N, j, n_{jk})$

$$n_{11} = a$$

$$(1, 1), (N, 1, a)i = 1$$

$$(2, 1), (N, 1, a)i = 2$$

$$n_{21} = c$$

$$(1, 1), (N, 2, c)i = 1$$

$$(2, 1), (N, 2, c)i = 2$$

$$n_{31} = e$$

$$(1, 1), (N, 3, e)i = 1$$

$$(2, 1), (N, 3, e)i = 2$$

.....

$$n_{32} = f$$

$$(1, 2), (N, 3, f)i = 1$$

$$(2, 2), (N, 3, f)i = 2$$

Matrix Matrix Multiplication (IV)

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} a & b \\ c & d \\ e & f \end{bmatrix} \longrightarrow \begin{bmatrix} 1a + 2c + 3e & 1b + 2d + 3f \\ 4a + 5c + 6e & 4b + 5d + 6f \end{bmatrix}$$

$((i, k), [(M, j, m_{ij}), (M, j, m_{ij}), \dots, (N, j, n_{jk}), (N, j, n_{jk}), \dots])$
 $(1, 1), [(M, 1, 1), (M, 2, 2), (M, 3, 3), (N, 1, a), (N, 2, c), (N, 3, e)]$
 $(1, 2), [(M, 1, 1), (M, 2, 2), (M, 3, 3), (N, 1, b), (N, 2, d), (N, 3, f)]$
 $(2, 1), [(M, 1, 4), (M, 2, 5), (M, 3, 6), (N, 1, a), (N, 2, c), (N, 3, e)]$
 $(2, 2), [(M, 1, 4), (M, 2, 5), (M, 3, 6), (N, 1, b), (N, 2, d), (N, 3, f)]$

Matrix Matrix Multiplication (IV)

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} a & b \\ c & d \\ e & f \end{bmatrix} \longrightarrow \begin{bmatrix} 1a + 2c + 3e & 1b + 2d + 3f \\ 4a + 5c + 6e & 4b + 5d + 6f \end{bmatrix}$$

$$[(M, 1, 1), (M, 2, 2), (M, 3, 3), (N, 1, a), (N, 2, c), (N, 3, e)]$$

$$list_M = [(M, 1, 1), (M, 2, 2), (M, 3, 3)]$$

$$list_N = [(N, 1, a), (N, 2, c), (N, 3, e)]$$

$$P(1, 1) = 1a + 2c + 3e$$

$$P(1, 2) = 1b + 2d + 3f$$

$$P(2, 1) = 4a + 5c + 6e$$

$$P(2, 2) = 4b + 5d + 6f$$